

w11L51: W11L1: Text Summarization - LexRank

Text Summarization

What is a summary?

A summary is a text that is produced from one or more texts, that contains a **significant portion of the information** in the original text(s), and that is **no longer than half** of the original text(s). (Hovy, 2008)

What is text summarization?

Text summarization is the process of **distilling** the most important information from a source (or sources) to produce an **abridged** version for a particular *user* or *task*. (Mani and Maybury, 2001)

Humans have an incredible capacity to condense information down to the **critical bit**.

- *Calvin Coolidge (heard lecture for 4 hours), on being asked what a clergyman preaching on sin said.*
 - *"He said he is against it."*

Automatic Text Summarization

Goal of a Text Summarization System

To give an **overview** of the original document in a **shorter** period of time.

Summarization Applications

- **outlines or abstracts** of any document, news article etc.
- **summaries** of email threads
- **action items** from a meeting
- **simplifying** text by compressing sentences

Application: Generating Snippets

Robert O'Neill taking credit for killing Osama bin Laden sparks debate

Hindustan Times - 1 hour ago

Some special operations service members and veterans are unhappy that one of their own has taken credit publicly for killing Osama bin Laden.

It's been special knock as wait has been long: Rayudu

Hindustan Times - 1 hour ago

An elated Ambati Rayudu said Friday that his maiden hundred in international cricket will certainly be a "special one" as it took a long time to come.

what is the relation between pressure and velocity

Web Images Videos News More Search tools

About 1,10,00,000 results (0.49 seconds)

fluid dynamics - Relation between pressure, velocity and ...
[physics.stackexchange.com/.../relation-between-pressure-velocity-and-ar...](#)
 In a nozzle, the exit **velocity** increases as per continuity equation as given by Bernoulli equation (incompressible fluid). **Pressure** is inversely proportional to ...

Chapter 9: Fluid Dynamics
[francesa.phy.cmich.edu/people/andy/physics110/book/.../Chapter9.htm](#)
 From practical experience we know that the velocity of fluid through the small ... we found a qualitative **relationship between pressure and velocity** in a fluid flow.

Bernoulli's Equation
[https://www.princeton.edu/~asmits/Bicycle_web/Bernoulli.html](#)
 ... can give great insight into the balance **between pressure, velocity** and elevation. ... When streamlines are parallel the **pressure** is constant across them, except ...

Pressure Vs velocity | Student Doctor Network
[forums.studentdoctor.net > ... > MCAT Study Question Q&A](#)
 Jul 21, 2009 - 8 posts - 3 authors
 Velocity increases with a decrease in pressure. Velocity... If you want to think of the **relationship between pressure and velocity**, you can use ...

Pressure - Velocity relation in fluids | CrazyEngineers
[www.crazyengineers.com > ... > Mechanical | Automobile | Aeronautics](#)
 May 31, 2009 - 20 posts - 6 authors
 The Darcys formula gives the **relation between** the Head Loss due to Friction in the flowing fluid wrt Flow **Velocity**, coefficient of friction, dia. of ...

you will get snippets

Automatic Text Summarization

Genres of Summary

Extract vs. Abstract <i>... lists fragments of text vs. re-phrases content coherently.</i> Single document vs. Multi-document <i>... based on one text vs. fuses together many texts.</i> Generic vs. Query-focused <i>... provides author's view vs. reflects user's interest.</i> <i>... like complex Q&A task</i>	Extract vs. Abstract <i>... lists fragments of text vs. re-phrases content coherently.</i> Single document vs. Multi-document <i>... based on one text vs. fuses together many texts.</i> Generic vs. Query-focused <i>... provides author's view vs. reflects user's interest.</i>
--	---

Will talk about Extractive for single document

Can use for Abstractive and multiple documents

Query-focused summarization can be thought of as a *complex question answering system*

Summarization: Main stages

1. Content Selection

- Choose sentences to extract from the document - find important sentences

2. Information Ordering

- Choose an order to place them in the summary
- how will i make it more readable presentable
- How it can be more coherent

3. Sentence realization

- Simplify the sentences

4. Removing Redundancy

- Increase **diversification** by removing redundant sentences

The most **basic algorithm** only does the first stage, **content selection**.

Unsupervised content selection; Luhn (1958)

Intuition

Choose **sentences** that have **salient** or **informative words** (what is the information content)

1. TF-IDF weight

Two approaches to define salient words

- *tf-idf*: weigh each word w_i in document j by tf-idf

$$weight(w_i) = tf_{ij} \times idf_i$$

2. What are domain specific words - take sentences which have **topic signature**

- *Topic signatures*: choose a smaller set of salient words, specific to that domain

$$weight(w_i) = 1 \text{ if } w_i \text{ is a specific term (use mutual information)}$$


Based on weights from one of above two approaches - take average

Weighing a sentence

$$weight(s) = \frac{1}{|S|} \sum_{w \in S} weight(w)$$


LexRank: A Graph-based approach

- Principled way of approach
 -
1. Document Representation



Assign tf-idf eg. 0.1, 0.15, 0.05 to each sentence

2. Finding the most salient sentences (Assignment question)



What is the cosine similarity between sentences?

Hypothesis: it is about underlying theme


- Sentences should be talking about **theme**
- This sentence should be similar to many other sentences
 - It will be important
 - Other important sentences should also be relevant

In Page rank - if many pages are linked to it


Sentence Centrality Measure

Finding the most salient sentences

1. A document graph is constructed with sentences as the vertices



2. A sentence similarity function is used to calculate the edge weights.



$$\tilde{M} = \begin{bmatrix} 0.0 & 0.5 & 0.0 & 0.4 & 0.1 \\ 0.5 & 0.0 & 0.5 & 0.0 & 0.0 \\ 0.0 & 0.5 & 0.0 & 0.5 & 0.0 \\ 0.4 & 0.0 & 0.4 & 0.0 & 0.2 \\ 0.3 & 0.0 & 0.0 & 0.7 & 0.0 \end{bmatrix}$$

N sentences: N x N pairs

3. PageRank based algorithm is used to compute the sentence centrality vector I .

Informative vectors

$$I_j = \mu \cdot \sum_{\forall k \neq j} I_k \cdot \tilde{M}_{k,j} + \frac{1 - \mu}{|S|}$$

$$I = \begin{bmatrix} 0.22 & 0.18 & 0.2 & 0.3 & 0.1 \end{bmatrix}$$


Node: that has highest weight is - most important

19min: S_{ij} - how similar two sentences are

Handwritten diagram showing a matrix A with elements S_{ij} . The matrix is labeled as a **row-stochastic matrix**. The text **PageRank** is written to the left. A dot product is indicated as $\cdot \rightarrow 1$.

Convert it to **row-stochastic matrix** - M -tilda

- All elements in individual rows sum to 1
- How does it help?
 - Gives probability - how can we see it as probability
 - Going from node i to node j




Doing random walk, you can go from node i to other nodes

Handwritten equation: $V = V \cdot M$


When you do random walk, it converges to V is page rank vector

A web page is important if many pages link to it

- How do they correlate?



you will stay at important node for longer time



As equation: importance of node

Importance of node = importance of other nodes * prob of linking from that node


Iteratively compute all I values

Start with v_0 , multiply by $\tilde{M}/v_1$ and so on

it will converge

where v_t and v_{t-1} are same

Keep on multiplying by M-tilda -> find the difference -> at sometime below threshold



$$I_j = \mu \cdot \sum_{\forall k \neq j} I_k \cdot \tilde{M}_{k,j} + \frac{1 - \mu}{|S|}$$


$$I = \begin{bmatrix} 0.22 & 0.18 & 0.2 & 0.3 & 0.1 \end{bmatrix}$$

Actual formula

What is

S - no. of sentences

mu is small teleport probability'



No out going node - random walk is stalled

- Assign small prob to all nodes - which is (1-mu)/S going to all node
- To balance have first term mu * M-tilda-ij

Mu is 0.85 - roughly taken

You will find page rank values $I = [0.22 \quad 0.18 \quad 0.2 \quad 0.3 \quad 0.1]$

- 4th sentence has highest prob (0.3)
- If you want to take longer take 1st (0.22 prob) or 3rd (0.2 prob)
- Take one of them as required

For diversity

- Give less weight to a sentence which is similar to other sentence

Removing Redundant Sentences

Maximal Marginal Relevance

- *How much more relevance are you bringing in*


- An iterative method for content selection from a selected list of important sentences

- Iteratively choose the best sentence to insert in the summary that is minimally redundant with the summary so far (Sum)

$$Inf(s)_{MMR} = \max_{s \in D} (Inf(s) - \lambda \cdot sim(s, Sum))$$


where $Inf(s)$ denotes the informativeness score of a sentence

- Minimally redundant with summary so far



Again rerank -> informative score - lambda similarity

Do for all sentences



Suppose it comes out to be
Sort and take one with highest - 0.4

Now you have two sentences

Again compute - take highest

- Until you get desired no. of sentences
- You have diversity

w11L52: W11L2: Optimization Based Approaches for Summarization

Discussed:

Graph based method - LexRank - take sentence which has high PageRank

Find diverse sentence - MMR - Max Margin Relevance algorithm

- Sort and chose sentence which has highest value
- Doesn't give flexibility in how you want to choose

You can modify it for different tasks: give flexibility to

- Define own optimization function
- Add constraints

Global Inference

$$D = t_1, t_2, \dots, t_{n-1}, t_n$$

Let us **define document D** with n textual units

Textual units are sentences - n sentences

- Let $Rel(i)$ be the relevance of t_i to be in the summary
- Let $Red(i, j)$ be the redundancy between t_i and t_j
- Let $l(i)$ be the length of t_i

Relevance:

- weights - do average
- Find page rank
- How many named entities are there
- What is the redundancy - we don't want multiple diverse sentences in summary
 - Cosine similarity
 - How much apart in the document - near, far apart
 - No. of words in text

Inference Problem



- The inference problem is to select a subset S of textual units from D such that summary score of S , i.e., $s(S)$, is maximized.
- $S = \arg \max_{S \subseteq D} \left[\sum_{t_i \in S} Rel(i) - \sum_{t_i, t_j \in S, i < j} Red(i, j) \right]$
such that $\sum_{t_i \in S} l(i) \leq K$, where k denotes the maximum length of the summary

Want summary of length k


- What should be optimisation criteria
- Select subset whose score is maximum
- [relevance - redundancy]
- Will not repeat pair so $i < j$
- Select subset which maximises it
- **Max rel, Min redundancy (Assignment question 1)**
- Constraint: K max length of summary
- How do you achieve this - maximise optimisation

A Greedy Solution

1. Sort D so that $Rel(i) > Rel(i + 1) \forall i$
2. $S = \{t_1\}$
3. while $\sum_{t_i \in S} l(i) < K$
4. $t_j = \arg \max_{t_j \in D - S} s(S \cup \{t_j\})$
5. $S = S \cup \{t_j\}$
6. return S




first sort t_1 to t_n , assume order is still t_1 to t_n




Take t_1 in summary.

Sum of rel t_1 and rel t_2 minus - (red of t_1 and t_2)

Add t_2, t_3



Again one by one add to Summary



Add one by one which one gives highest score

1. Sort D so that $Rel(i) > Rel(i + 1) \forall i$
2. $S = \{t_1\}$
3. while $\sum_{t_i \in S} l(i) < K$
4. $t_j = \arg \max_{t_j \in D - S} s(S \cup \{t_j\})$
5. $S = S \cup \{t_j\}$
6. return S

3 - until desired length

4 - take remaining sentences - add one which has highest score

Achieved both relevance and diversity

- It is greedy approach
- Not going to an optimal solution - near optimal
- Good enough for decent solution

You can use Dynamic Programming

- Storing results, what is the best way of achieving summary 5, 10, 15

Integer Linear Programming (ILP)

- Greedy algorithm (previously) is an approximate solution
- Use **exact** solution algorithms with ILP - optimal value
- ILP is a **constrained optimization problem**
 - E.g. Summation less than K
- Many solvers on the web - you can take any one of them
- To understand
 - Define the **constraints** based on **relevance** and **redundancy** for summarization

Sentence Level ILP Formulation

Optimization Function

$$\text{maximize } \sum_i \alpha_i \text{Rel}(i) - \sum_{i < j} \alpha_{ij} \text{Red}(i, j)$$

We are not taking subsets

Handwritten equation: $\sum_i \alpha_i \text{Rel}(i) - \sum_{i,j} \alpha_{ij} \text{Red}(i,j)$. Arrows point from the i in the first sum to $\{0,1\}$ and from the i,j in the second sum to $\{0,1\}^2$.

alpha-i : whether ith unit is included in summary or not - should have value of 0 or 1

- 0 not in summary, 1 in summary

alpha-ij : can take value 0 or 1

- ij both are in summary 1

Handwritten equation: $\alpha_i = 1 \forall i$ and $\alpha_{ij} = 0 \forall i,j$.

Can put solution, you are taking whole document

Constraints

Handwritten constraint equation: $\sum \alpha_i l(i) \leq K$.

Length of summary less than K -

Handwritten equation: $\alpha_i = 1 \forall i$, circled and crossed out with a large X.

Will ensure this solution is not acceptable

Whenever α_i to be 1 and $\alpha_{ij} = 0$ - don't want

- Also don't want $\alpha_{ij} = 1$ and α_i to be 0

When α_i and α_j are 1, want α_{ij} has to be 1

- Need some more constraints

How do we put things in constraint

such that $\forall i, j$:

- $\alpha_i, \alpha_{ij} \in \{0, 1\}$
- $\sum_i \alpha_i l(i) \leq K$
- $\alpha_{ij} - \alpha_i \leq 0$
- $\alpha_{ij} - \alpha_j \leq 0$
- $\alpha_i + \alpha_j - \alpha_{ij} \leq 1$

Handwritten notes showing constraints and variable values. It includes the constraint $l(i) \leq K$, and a table of values for α_{ij} and α_i, α_j . A circled note says $\alpha_{ij} = 1, \alpha_i = 1, \alpha_j = 1$.

Handwritten equation $\alpha_{ij} - \alpha_j \leq 0$.

understand same way

3rd constraint:

Handwritten equation $\alpha_i + \alpha_j - \alpha_{ij} \leq 1$.

1 or 1 or 0 ≤ 1


Can't have $\alpha_{ij} = 0$ when α_i and α_j are one

Now putting constraints in terms of Integer Linear Programming

Is generic enough

Depending on your task, you can define your own optimization function and constraints.

Doing summary from a research article:



- Has sections
- In summary want at least one sentence from each section
 - Define section information

News article - Data is temporal

- If two information are in two time point, take information from later time point than previous time point - certain time period more than other time period
- Optimization fn - put weight (importance) to different days
- Max multiple things together
- Still be able to find solution using ILP

2. The next steps: Sentence Ordering

- Heuristic approach

Chronological ordering: the simplest method

1. List the sentences in the order, they appear in the document
- First sentence you will put first

Coherence

2. Choose orderings that make neighboring sentences **similar** (by cosine)
3. Choose orderings in which neighboring sentences discuss the **same entity, events**

Topical ordering

4. Learn the ordering of **topics** in the source documents

3. The next steps: Simplifying Sentences

- Can get more space

Parse sentences, use rules to decide which modifiers to **prune**

- **Initial adverbials:** *For example, on the other hand, as a matter of fact, at this point, ...*
 - *Not important to convey information*
- **PPs without named entities:**
 - The commercial fishing restrictions in Washington will not be lifted unless the salmon population increases [*PP to a sustainable number*] *Not important to convey information*
- **Attribution clauses:**
 - Rebels agreed to talks with government officials, *international observers said Tuesday*
- **Appositives:**
 - Rajan, 28, *an artist who was living at the time in Philadelphia*, found the inspiration in the back of city magazines

These rules can be given to the system, or learn the rules (by inputs provided by humans)

w11L53: W11L3: Summarization: Evaluation (20min)

System Evaluation

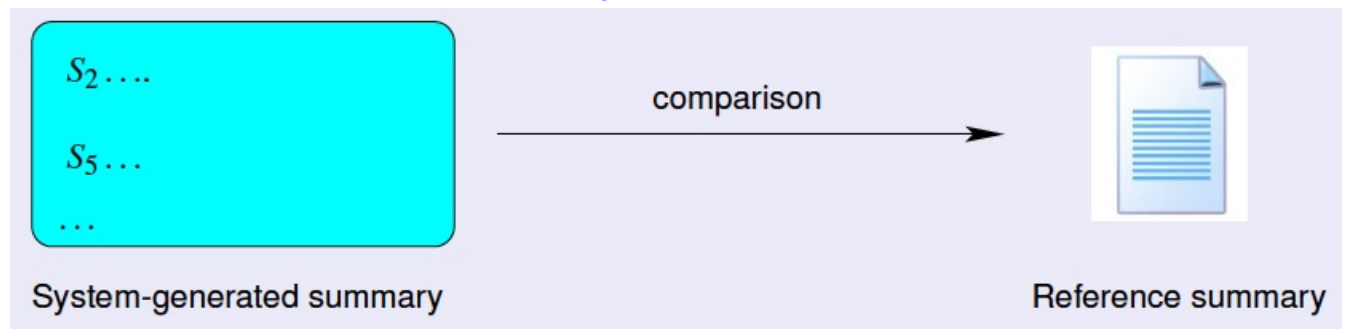
- Human should say what is the good information coverage (on scale 1-5)
- What kind of diversity (on scale 1-5)
- Readable - coherent (on scale 1-5)

Can you do it automatically?

- Compare system generated summary with Reference summary

Evaluation Criteria

System Evaluation



ROUGE


Recall Oriented Understudy for Gisting Evaluation

- *Not as good as human evaluation but much more **convenient***
- *Human evaluation may vary*

Toolkit available for download.

DUC (o2, o3 datasets): Document Understanding Conferences

- 400 documents, for each document -> provide manual summary



vary the algorithm and compare

ROUGE for evaluation (has many variations - unigram, bigram, longest common sub-sentence)
Popular - ROGUE-1 (unigram) ROGUE-2 (bigram)

Given a document D , and an automatic summary X :

- Have N humans produce a set of reference summaries of D ($N \geq 1$)
- Run system, giving automatic summary X
- What percentage of the n-grams from the reference summaries appear in X ?

$$ROUGE - 2 = \frac{\sum_{S \in \{RefSums\}} \sum_{bi-gram \in S} Count_{match}(bi-gram)}{\sum_{S \in \{RefSums\}} \sum_{bi-gram \in S} Count(bi-gram)}$$

It is like RECALL:

For all ref summaries - how many bigrams in system generated summary / by

- How many bigrams are present in system summary

If bigram appears in system generated - ROUGE will be higher

ROUGE Example

Reference Summaries

Human 1: water spinach is a green leafy vegetable grown in the tropics.

bigram = no. of words - 1 = 10 bigrams

Human 2: water spinach is a semi-aquatic tropical plant grown as a vegetable.

Human 3: water spinach is a commonly eaten leaf vegetable of Asia

System Summary

water spinach is a leaf vegetable commonly eaten in tropical areas of Asia.

Denominator:

How many bigrams are common with summary (bigram = no. of words - 1) = 10+10+9

Numerator:

Common bigrams in human-1 + Common bigrams in human-2 + Common bigrams in human-3

ROUGE-2 score $\frac{3 + 3 + 6}{10 + 10 + 9} = 12/29 = 0.413$

Computer for all document in your corpus

Toolkit available for download.

- It can compute ROUGE score and confidence interval
- Can know how much scores are different

I want summary of length

- Stop at 105 or 107 instead of just length upto 100 (1st 100 words can be taken)
- I want to use stemming
 - Boy and boys will start matching
 - I will remove stop word (is and a)

Define manually some semantic content


- should have max coverage

Can you find some property if it is a good summary

We focused on Generic and single document summarisation

Further Discussions



- Multi-document summarization
 - Can apply same algorithm - LexRank
 - You want to give different weightage to different documents
 - Different way of similarity for
 - Statements with in documents and across documents
 - Overall idea is same
- Query-specific summarization
 - Find out sentences similar to query, then do summarisation
 - Do larger summary than you wanted
 - Then find out which ones are more closer to the query
 - Can do query specific multi document query
 - First select documents, then sentences which are similar
- Abstractive summarization
 - Slightly trickier
 - You are fusing information
 - Apply extractive summary, get a larger set (say summary of length 200)



s1, s2 and s3 are similar, fuse them

Words are common or some information provided by each sentence

Construct a graph



Wherever diverse use generative methods

- Paths can be based on unigram, bigram

Applications:

- Scientific articles
- Twitter stream
- Quora, stackover flow answers
- Disaster events

w11L54: W11L4: Text Classification - I

Example: Positive or negative movie review?



- unbelievably disappointing



- Full of zany characters and richly applied satire, and some great plot twists



- this is the greatest screwball comedy ever filmed



- It was pathetic. The worst part about it was the boxing scenes.

Example:

Male or Female Author? What is the age of author (20s, 30s, old age)

Which country, who is the author?

Authors will have some trait

1. By 1925 present-day Vietnam was divided into three parts under French colonial rule. The southern region embracing Saigon and the Mekong delta was the colony of Cochin-China; the central area with its imperial capital at Hue was the protectorate of Annam...
2. Clara never failed to be astonished by the extraordinary felicity of her own name. She found it hard to trust herself to the mercy of fate, which had managed over the years to convert her greatest shame into one of her greatest assets...

Male - write about different fact

Female - opinion and subjective information

Example: What is the subject of this article? Or Categories of Article

MEDLINE Article



MeSH Subject Category Hierarchy

- Antagonists and Inhibitors
- Blood Supply
- Chemistry
- Drug Therapy
- Embryology
- Epidemiology
- ...

Text Classification

- Assigning subject categories, topics, or genres - category of text classification
- Spam detection
 - Email, Tweet, Movie review, Product review - is it spam or not
- Authorship identification
- Age/gender identification
- Language identification
 - English vs Hindi - Hindi in Devanagari - use roman script
 - They have same script - need to go to word level
 - Same translation - pure - english. pure (poore - whole in hindi)
- Sentiment analysis
 -

Text classification: problem definition

d = piece of text, Each document belongs to one class

Input

- A document d
- A fixed set of classes $C = \{c_1, c_2, \dots, c_n\}$

Output

A predicted class $c \in C$

1. Simplest method:

Classification Methods: Hand-coded rules

- Rules based on combinations of words or other features

Spam

black-list-address OR ("dollars" AND "have been selected")

Pros and Cons

Accuracy can be high if rules carefully refined by expert, *but building and maintaining these rules is expensive.*

Need to have sufficient rule, but will not know how many

Will not have high recall

2. Classification Methods: Supervised Machine Learning

- Naive Bayes
- Logistic regression (LR) - similar to MaxEnt - what is the probability of a class given input x
 - Linear function over various features (exponent of lambda - si)
- Support-vector machines (SVM)
- Decision Tree (DT, RF, XGBoost) ...


Naive Bayes - Powerful baseline

Naive Bayes Intuition

- Simple classification method based on Bayes' rule
- Relies on very simple representation of document - Bag of words

Bag of words for document classification

- Set of words. Order of the words don't matter



parser
language
label
translation
...

?

Machine Learning	NLP	Garbage Collection	Planning	GUI
learning	<u>parser</u>	garbage	planning	...
<u>training</u>	tag	collection	temporal	
algorithm	training	memory	reasoning	
shrinkage	<u>translation</u>	optimization	plan	
network...	<u>language...</u>	region...	<u>language...</u>	

Document = bag of words Document is assigned class

At test time (left) you need to find class of document

Bayes' rule for documents and classes

For a document d and a class c $P(c d) = \frac{P(d c)P(c)}{P(d)}$ Find prob of all different classes Find which one has max prob	Naive Bayes Classifier $c_{MAP} = \arg \max_{c \in C} P(c d)$ $= \arg \max_{c \in C} P(d c)P(c)$ $= \arg \max_{c \in C} P(x_1, x_2, \dots, x_n c)P(c)$
--	--

How do you compute $P(c|d)$

- Generative - have class then generate the document
- Do you want to generate +ve or -ve review

$$P(c|d) = \frac{P(d|c) P(c)}{P(d)}$$

class document

$$C_{MMP} = \operatorname{argmax}_{c \in C} P(c|d)$$

C → d_i

Hence can't compute $P(c|d)$ directly

Use Bayes rule

$$C_{MMP} = \operatorname{argmax}_{c \in C} \frac{P(c|d)}{P(d)}$$
$$= \operatorname{argmax}_c \frac{P(d|c) P(c)}{P(d)}$$

$P(d)$ is common for all classes

As $P(d)$ is common across all classes we can remove $P(d)$ from denominator

$$\operatorname{argmax}_c \frac{P(d|c) P(c)}{P(d)}$$

features


$$\operatorname{argmax}_c P(x_1, \dots, x_n | c) P(c)$$

$P(c)$ - prior prob of class - how often class c occurs

$P(d|c)$ - given data what is the probability

- What is your document - bag of words x_1 ,
- X can be words, time of document etc.

What is the probability of getting all features



Don't worry about join probability !!!

$P(x_1, x_2, \dots, x_n | c)$ Naive Bayes classification assumptions

1. Bag of words assumption

Assume that the **position** of a word in the document doesn't matter

2. Conditional Independence

Assume the feature probabilities $P(x_i | c_j)$ are independent given the class c_j .

$$P(x_1, x_2, \dots, x_n | c) = P(x_1 | c) \cdot P(x_2 | c) \dots P(x_n | c)$$

$$C_{NB} = \arg \max_{c \in C} P(c) \prod_{x \in X} P(x | c)$$

X - different words

How do you compute probabilities from training data?

Learning the model parameters

Maximum Likelihood Estimate - $P(c)$ what is the probability of this class in data

$$\hat{P}(c_j) = \frac{\text{doc-count}(C = c_j)}{N_{\text{doc}}}$$

How many documents have been labelled with this class / no. of documents in training data

$$\hat{P}(w_i | c_j) = \frac{\text{count}(w_i, c_j)}{\sum_{w \in V} \text{count}(w, c_j)}$$

Prob of word given class = no. of time this words occurs in c_j / count of all words in the class c_j


- Different tokens

Problem with MLE

Suppose in the training data, we haven't seen the word "*fantastic*", classified in the topic 'positive'.

$$\hat{P}(\text{fantastic} | \text{positive}) = 0$$

$P(w_i | c_j)$



Only one word will make the prob = 0

$$c_{NB} = \arg \max_c \hat{P}(c) \prod_{x \in X} \hat{P}(x_i | c)$$

$$\frac{C(w_i, c_j)}{\sum_w C(w, c_j)} = p(w_i | c_j)$$

MLE

$$\frac{C(w_i, c_j) + 1}{\sum_w C(w, c_j) + V} = p(w_i | c_j)$$

MLE

size of vocab.

+1 in numerator

+V in denominator

Laplace (add-1) smoothing

$$\hat{P}(w_i | c) = \frac{\text{count}(w_i, c) + 1}{\sum_{w \in V} (\text{count}(w, c) + 1)}$$

$$= \frac{\text{count}(w_i, c) + 1}{(\sum_{w \in V} \text{count}(w, c)) + |V|}$$

So whole thing will not be zero

w11L55: W11L5: Text Classification - II

A worked example

$$\hat{P}(c) = \frac{N_c}{N}$$

$$\hat{P}(w|c) = \frac{\text{count}(w,c)+1}{\text{count}(c)+|V|}$$

	Doc	Words	Class
Training	1	Chinese Beijing Chinese	c
	2	Chinese Chinese Shanghai	c
	3	Chinese Macao	c
	4	Tokyo Japan Chinese	j
Test	5	Chinese Chinese Chinese Tokyo Japan	?

Class c occurs 3 times out of 4

$$\hat{P}(c) = \frac{3}{4}$$

$$\hat{P}(j) = \frac{1}{4}$$

Want to assign test to some class:

Test Chinese-3, Tokyo-1, Japan-1

$$\frac{\hat{P}(c) \cdot P(\text{Chinese}|c)^3 P(\text{Tokyo}|c) P(\text{Japan}|c)}{\hat{P}(j) \cdot P(\text{Chinese}|j)^3 P(\text{Tokyo}|j) P(\text{Japan}|j)}$$

Test Chinese-3, Tokyo-1, Japan-1

How to compute prob $P(\text{chinese}|c) =$

$$\hat{P}(c) = \frac{3}{4}$$

$$P(\text{Chinese}|c) = \frac{c(\text{Chinese},c)+1}{n(c)+|V|}$$

$$\hat{P}(j) = \frac{1}{4}$$

$$P(\text{Chinese}|j) = \frac{c(\text{Chinese},j)+1}{n(j)+|V|}$$

Chinese occurs in class c -5 How many different tokens - 8

$$P(\text{Chinese}|c) = \frac{c(\text{Chinese},c)+1}{n(c)+|V|} = \frac{5+1}{8+6} = \frac{6}{14}$$

$$P(\text{Chinese}|j) = \frac{c(\text{Chinese},j)+1}{n(j)+|V|} = \frac{1+1}{3+6} = \frac{2}{9}$$

$P(\text{chinese} | c) = 3/7$ $P(\text{chinese} | j) = 2/9$

Computer $P(\text{tokyo} | c) = 1/14$ and $P(\text{japan} | c) = 1/14$ $P(\text{tokyo} | j) = 2/9$ and $P(\text{japan} | j) = 2/9$

Compute $P(c | d5)$

Compute $P(j | d5)$

Handwritten notes on a whiteboard:

$$\hat{P}(c) = \frac{P(\text{chinese}|c)^3 P(\text{Tokyo}|c) P(\text{Japan}|c)}{P(c)}$$

$$\hat{P}(j) = \frac{P(\text{chinese}|j)^3 P(\text{Tokyo}|j) P(\text{Japan}|j)}{P(j)}$$

Test: Chinese-3, Tokyo-1, Japan-1

$$P(c | d5) \propto \frac{3}{4} * \left(\frac{3}{7}\right)^3 * \frac{1}{14} * \frac{1}{14} \approx 0.0003$$

$$P(j | d5) \propto \frac{1}{4} * \left(\frac{2}{9}\right)^3 * \frac{2}{9} * \frac{2}{9} \approx 0.0001$$

Class C has higher prob than class j

Priors:

$$P(c) = \frac{3}{4}$$

$$P(j) = \frac{1}{4}$$

Choosing a class:

$$P(c | d5) \propto \frac{3}{4} * \left(\frac{3}{7}\right)^3 * \frac{1}{14} * \frac{1}{14} \approx 0.0003$$

Conditional Probabilities:

$$P(\text{Chinese} | c) = \frac{(5+1)}{(8+6)} = \frac{6}{14} = \frac{3}{7}$$

$$P(\text{Tokyo} | c) = \frac{(0+1)}{(8+6)} = \frac{1}{14}$$

$$P(\text{Japan} | c) = \frac{(0+1)}{(8+6)} = \frac{1}{14}$$

$$P(\text{Chinese} | j) = \frac{(1+1)}{(3+6)} = \frac{2}{9}$$

$$P(\text{Tokyo} | j) = \frac{(1+1)}{(3+6)} = \frac{2}{9}$$

$$P(\text{Japan} | j) = \frac{(1+1)}{(3+6)} = \frac{2}{9}$$

$$P(j | d5) \propto \frac{1}{4} * \left(\frac{2}{9}\right)^3 * \frac{2}{9} * \frac{2}{9} \approx 0.0001$$

Naive Bayes and Language Modeling

In general, NB classifier can use any feature

- URL, email addresses, dictionaries, network features (e.g. email from this person)

*But if we use only the **word** features and all the words in the text*

- Naive Bayes has an important **similarity** to language modeling.
- Each **class** can be thought of as a separate **unigram** language model.
- For each class you are building Language Model
- Which language model assigns higher probability

Naive Bayes as Language Modeling

- For +ve class - build a unigram model
- For -ve class - build a unigram model

Which class assigns a higher probability to the sentence?

Model pos		Model neg						
0.1	i	0.2	i	i	love	this	fun	film
0.1	love	0.001	love					
0.01	this	0.01	this					
0.05	fun	0.005	fun					
0.1	film	0.1	film					

i	love	this	fun	film
0.1 0.2	0.1 0.001	0.01 0.01	0.05 0.005	0.1 0.1

$P(s|\text{pos}) > P(s|\text{neg})$

At test time you see word : i love this fun film. +ve model is giving higher probability
 Another way of thinking about this problem

Naive Bayes: More than Two Classes

Multi-value classification


- Categories are not mutually exclusive

Multi-value classification

A document can belong to 0, 1 or > 1 classes

Make different binary classifier for each class

C_1 is positive example, $\neg C_1$ is negative



Each document at test time, might be given multiple labels

Handling Multi-value classification

- For each class $c \in C$, build a classifier γ_c to distinguish c from all other classes $c' \in C$
- Given test-doc d , evaluate it for membership in each class using each γ_c
- d belongs to any class for which γ_c returns true

Multi-class classification using Binary Classification

One-of or multinomial classification

Classes are mutually exclusive: each document in exactly one class

Binary classifiers may also be used

- For each class $c \in C$, build a classifier γ_c to distinguish c from all other classes $c' \in C$
- Given test-doc d , evaluate it for membership in each class using each γ_c
- d belongs to one class with maximum score

What prob the classifier gave -> take the max probability -> assign that class

Evaluation:

The diagram shows a vertical list of test data points $d_1, \dots, d_n$ on the left. Each point d_i is connected by a horizontal line to a predicted class $c_1, c_3, c_4, \dots, c_k$ in the middle column. To the right of each predicted class is the true class $c_1, c_3, c_2, \dots, c_4$. The labels 'Predicted class' and 'True class' are written above their respective columns.

Evaluation measures

Accuracy: What fraction of time same answer

Evaluation: Constructing Confusion matrix c

For each pair of classes $\langle c_1, c_2 \rangle$ how many documents from c_1 were incorrectly assigned to c_2 ? (when $c_2 \neq c_1$)

Docs in test set	Assigned UK	Assigned poultry	Assigned wheat	Assigned coffee	Assigned interest	Assigned trade
True UK	95	1	13	0	1	0
True poultry	0	1	0	0	0	0
True wheat	10	90	0	1	0	0
True coffee	0	0	0	34	3	7
True interest	-	1	2	13	26	5
True trade	0	0	2	14	5	10

95 = assigned 95 they were true - classifier confused between UK and wheat

90 = poultry to wheat - really confused

How many assigned UK 95+10 (col)

True UK = 110 (row)

Precision: of all assigned UK, 95 were correct

Recall: of UK class - recalled 95 out of 110

Precision Micro	Precision Macro
-----------------	-----------------

Per class evaluation measures

Recall

Fraction of docs in class i classified correctly: $\frac{c_{ii}}{\sum_j c_{ij}}$

Cii / Add row-ij

Precision

Fraction of docs assigned class i that are actually about class i : $\frac{c_{ii}}{\sum_i c_{ji}}$

cii/ add col-ji

Accuracy

Fraction of docs classified correctly: $\frac{\sum_i c_{ii}}{N}$

Add diagonal element / total number of docs

This is for individual classes

Micro- vs. Macro-Average

If we have more than one class, how do we **combine multiple performance measures** into **one quantity**?

Macro-averaging

Compute performance (precision, accuracy) for each class, then **average** $((0.7 + 0.9) / 2)$

Micro-averaging

Collect decisions for all the classes, compute contingency table, evaluate.

Micro- vs. Macro-Average

Class 1

	Truth: yes	Truth: no
Classifier: yes	10	10
Classifier: no	10	970

Class 2

	Truth: yes	Truth: no
Classifier: yes	90	10
Classifier: no	10	890

Micro Ave. Table

	Truth: yes	Truth: no
Classifier: yes	100	20
Classifier: no	20	1860

Precision for class1: $10/20 = 0.5$ Precision for class2 $90/100 = 0.9$

Macro-averaged precision: $(0.5 + 0.9)/2 = 0.7$

Given equal weight

Compute on 3rd table:

Micro-averaged precision: $100/120 = 0.83$

Biased towards having more no. of instances

Micro-averaged score is dominated by score on common classes

Use Macro - for imbalance dataset

Resample:

- Undersample
- Oversample (from minority class)
- Or do both
-

